

Week 8: Homework 2: Project: Movie Recommendation with MLlib - Collaborative Filtering (implementation 3)

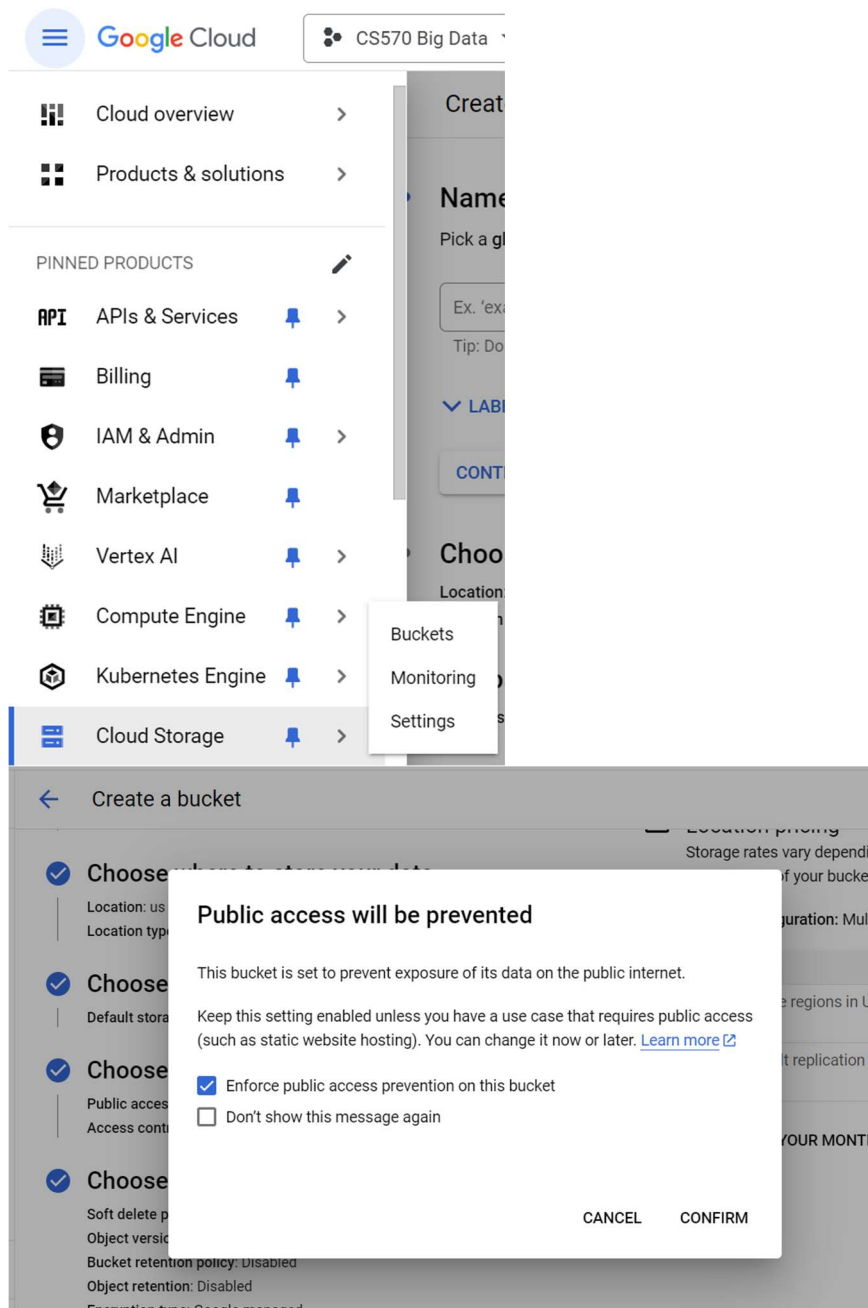
Step-by-Step Guide for Deployment on GCP with Correct File Paths

1. Upload Data and Scripts to GCS

Ensure that `movies.csv`, `ratings.csv`, and your PySpark script (e.g., `Recommendation_Engine_MovieLens.py`) are already uploaded to your GCS bucket.

2. Create a Google Cloud Storage (GCS) Bucket

Create a bucket in GCS to store your scripts and data:



3. Upload Data and Scripts to GCS

Upload the `movies.csv`, `ratings.csv`, and your PySpark script to your GCS bucket:

Bucket details

Location: us (multiple regions in United States) | Storage class: Standard | Public access: Not public | Protection: Soft Delete

OBJECTS | CONFIGURATION | PERMISSIONS | PROTECTION | LIFECYCLE | OBSERVABILITY | INVENTORY REPORTS

Folder browser: movie_recommendation_with_mllib_collaborative_filte

Actions: UPLOAD FILES, UPLOAD FOLDER, CREATE FOLDER, TRANSFER DATA, MANAGE HOLDS, EDIT RETENTION, DOWNLOAD, DELETE

Name	Size	Type	Created
movies.csv	482.8 KB	text/csv	Jul 17, 2024, 1:32
ratings.csv	2.4 MB	text/csv	Jul 17, 2024, 1:32
tags.csv	115.9 KB	text/csv	Jul 17, 2024, 1:32

4. Modify the PySpark Script to Use GCS Paths

Update your PySpark script to read the files from GCS. Use command-line arguments to pass the paths of the CSV files, enhancing the script's flexibility. Upload the modified script to the bucket:

```
shagos90499@cloudshell:~ (cs570-big-data-424809) $ vim Recommendation_Engine_MovieLens.py
shagos90499@cloudshell:~ (cs570-big-data-424809) $
```

Here is the script:

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, explode
from pyspark.ml.evaluation import RegressionEvaluator
from pyspark.ml.recommendation import ALS
from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
import argparse

# Parse command-line arguments
parser = argparse.ArgumentParser()
parser.add_argument('--input_path_movies', required=True)
```

```

parser.add_argument('--input_path_ratings', required=True)
args = parser.parse_args()

# Initialize Spark session
spark = SparkSession.builder.appName('Recommendations').getOrCreate()

# Load data from GCS
movies = spark.read.csv(args.input_path_movies, header=True)
ratings = spark.read.csv(args.input_path_ratings, header=True)

# Preprocess data
ratings = ratings \
    .withColumn('userId', col('userId').cast('integer')) \
    .withColumn('movieId', col('movieId').cast('integer')) \
    .withColumn('rating', col('rating').cast('float')) \
    .drop('timestamp')

# Split data into training and testing sets
(train, test) = ratings.randomSplit([0.8, 0.2], seed=1234)

# Build ALS model
als = ALS(userCol="userId", itemCol="movieId", ratingCol="rating",
nonnegative=True, implicitPrefs=False, coldStartStrategy="drop")
param_grid = ParamGridBuilder() \
    .addGrid(als.rank, [10, 50, 100, 150]) \
    .addGrid(als.regParam, [.01, .05, .1, .15]) \
    .build()
evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",
predictionCol="prediction")
cv = CrossValidator(estimator=als, estimatorParamMaps=param_grid,
evaluator=evaluator, numFolds=5)

# Train model
model = cv.fit(train)
best_model = model.bestModel

# Evaluate model
test_predictions = best_model.transform(test)
RMSE = evaluator.evaluate(test_predictions)
print(f"Root-mean-square error = {RMSE}")

# Generate recommendations
nrecommendations = best_model.recommendForAllUsers(10)
nrecommendations = nrecommendations \
    .withColumn("rec_exp", explode("recommendations")) \
    .select('userId', col("rec_exp.movieId"), col("rec_exp.rating"))
nrecommendations.show()

# Join with movie titles for better interpretability
nrecommendations.join(movies, on='movieId').filter('userId = 100').show()
ratings.join(movies, on='movieId').filter('userId = 100').sort('rating',
ascending=False).limit(10).show()

# Stop Spark session
spark.stop()

```

```

# Build ALS model
als = ALS(userCol="userId", itemCol="movieId", ratingCol="rating", nonnegative=True, implicitPrefs=False, cold_start_threshold=10, param_grid = ParamGridBuilder() \
# Split data into training and testing sets
(train, test) = ratings.randomSplit([0.8, 0.2], seed=1234)

# Build ALS model
als = ALS(userCol="userId", itemCol="movieId", ratingCol="rating", nonnegative=True, implicitPrefs=False, cold_start_threshold=10, param_grid = ParamGridBuilder() \
.addGrid(als.rank, [10, 50, 100, 150]) \
.addGrid(als.regParam, [.01, .05, .1, .15]) \
.build()
evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating", predictionCol="prediction")
cv = CrossValidator(estimator=als, estimatorParamMaps=param_grid, evaluator=evaluator, numFolds=5)

# Train model
model = cv.fit(train)
best_model = model.bestModel

# Evaluate model
test_predictions = best_model.transform(test)
RMSE = evaluator.evaluate(test_predictions)
print(f"Root-mean-square error = {RMSE}")

# Generate recommendations
nrecommendations = best_model.recommendForAllUsers(10)
nrecommendations = nrecommendations \
.withColumn("rec_exp", explode("recommendations")) \
.select('userId', col("rec_exp.movieId"), col("rec_exp.rating"))
nrecommendations.show()

# Join with movie titles for better interpretability
nrecommendations.join(movies, on='movieId').filter('userId = 100').show()
ratings.join(movies, on='movieId').filter('userId = 100').sort('rating', ascending=False).limit(10).show()

```

Upload the file:

```

gsutil cp Recommendation_Engine_MovieLens.py
gs://movie_recommendation_with_mllib_collaborative_filte

```

```

shagos90499@cloudshell:~ (cs570-big-data-424809) $ gsutil cp Recommendation_Engine_MovieLens.py gs://movie_recommendation_with_mllib_collaborative_filte
Copying file://Recommendation_Engine_MovieLens.py [Content-Type=text/x-python]...
/ [1 files] 2.2 KiB/ 2.2 KiB
Operation completed over 1 objects/2.2 KiB.
shagos90499@cloudshell:~ (cs570-big-data-424809) $

```

5. Create the Cluster with the Desired Configuration

Create a Dataproc cluster:

```

gcloud dataproc clusters create spark-cluster \
--region us-west1 \
--zone us-west1-a \
--master-machine-type n1-standard-4 \
--worker-machine-type n1-standard-4 \
--num-workers 2

```

```
shagos90499@cloudshell:~ (cs570-big-data-424809)$ gcloud dataproc clusters create spark-cluster \
  --region us-west1 \
  --zone us-west1-a \
  --master-machine-type n1-standard-4 \
  --worker-machine-type n1-standard-4 \
  --num-workers 2
Waiting on operation [projects/cs570-big-data-424809/regions/us-west1/operations/65ebb5fa-f4e4-3975-958f-82a592bed5cf].
Waiting for cluster creation operation...
WARNING: No image specified. Using the default image version. It is recommended to select a specific image version in production, as the default image version may change at any time.
WARNING: Consider using Auto Zone rather than selecting a zone manually. See https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/auto-zone
WARNING: Failed to validate permissions required for default service account: '720083396959-compute@developer.gserviceaccount.com'. Cluster creation could still be successful if required permissions have been granted to the respective service accounts as mentioned in the document https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/service-accounts#dataproc_service_accounts_2. This could be due to Cloud Resource Manager API hasn't been enabled in your project '720083396959' before or it is disabled. Enable it by visiting 'https://console.developers.google.com/apis/api/cloudresourcemanager.googleapis.com/overview?project=720083396959'.
WARNING: The firewall rules for specified network or subnet would allow ingress traffic from 0.0.0.0/0, which could be a security risk.
WARNING: The specified custom staging bucket 'dataproc-staging-us-west1-720083396959-ncaxf7jl' is not using uniform bucket level access IAM configuration. It is recommended to update bucket to enable the same. See https://cloud.google.com/storage/docs/uniform-bucket-level-access.
Waiting for cluster creation operation...done.
Created [https://dataproc.googleapis.com/v1/projects/cs570-big-data-424809/regions/us-west1/clusters/spark-cluster] Cluster placed in zone [us-west1-a].
shagos90499@cloudshell:~ (cs570-big-data-424809)$
```

6. Submit the PySpark Job with GCS Paths

- Submit your PySpark job to the Dataproc cluster, specifying the GCS paths for the input files:

```
gcloud dataproc jobs submit pyspark
gs://movie_recommendation_with_mllib_collaborative_filte/Recommendation_Engine_MovieLens.py \
  --cluster=spark-cluster \
  --region=us-west1 \
  -- \
  --input_path_movies=gs://movie_recommendation_with_mllib_collaborative_filte/movies.csv \
  --input_path_ratings=gs://movie_recommendation_with_mllib_collaborative_filte/ratings.csv
```

Replace your-bucket-name with the actual name of your GCS bucket. By following these steps, your PySpark script will correctly read the files from GCS when running on GCP Dataproc.

- Test result: root mean square error is calculated.

```
shagos90499@cloudshell:~ (cs570-big-data-424809)$ gcloud dataproc jobs submit pyspark gs://movie_recommendation_with_mllib_collaborative_filte/Recommendation_Engine_MovieLens.py \
  --cluster=spark-cluster --region=us-west1 -- --input_path_movies=gs://movie_recommendation_with_mllib_collaborative_filte/movies.csv \
  --input_path_ratings=gs://movie_recommendation_with_mllib_collaborative_filte/ratings.csv
Job [e4151b89e5d24303a13656f17a345a38] submitted.
Waiting for job output...
24/07/17 09:00:51 INFO org.apache.spark.SparkEnv: Registering MapOutputTracker
24/07/17 09:00:51 INFO org.apache.spark.SparkEnv: Registering BlockManagerMaster
24/07/17 09:00:51 INFO org.apache.spark.SparkEnv: Registering BlockManagerMasterHeartbeat
24/07/17 09:00:51 INFO org.apache.spark.SparkEnv: Registering OutputCommitCoordinator
24/07/17 09:00:52 INFO org.sparkproject.jetty.util.log: Logging initialized @3709ms to org.sparkproject.jetty.util.log.Slf4jLog
24/07/17 09:00:52 INFO org.sparkproject.jetty.server.Server: jetty-9.4.40.v20210413; built: 2021-04-13T20:42:42.668Z; git: b881a572662e1943a14ae12e7e1207989f2174; jvm 1.8.0_412-b08
24/07/17 09:00:52 INFO org.sparkproject.jetty.server.Server: Started @3814ms
24/07/17 09:00:52 INFO org.sparkproject.jetty.server.AbstractConnector: Started ServerConnector@34be357[HTTP/1.1, (http/1.1)]{0.0.0.0:46581}
24/07/17 09:00:52 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at spark-cluster-m/10.138.0.28:8032
24/07/17 09:00:53 INFO org.apache.hadoop.yarn.client.AHSProxy: Connecting to Application History server at spark-cluster-m/10.138.0.28:10200
24/07/17 09:00:53 INFO org.apache.hadoop.conf.Configuration: resource-types.xml not found
24/07/17 09:00:54 INFO org.apache.hadoop.yarn.util.resource.ResourceUtils: Unable to find 'resource-types.xml'.
24/07/17 09:00:54 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl: Submitted application 1721206458042_0002
24/07/17 09:00:56 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at spark-cluster-m/10.138.0.28:8030
24/07/17 09:00:58 INFO com.google.cloud.hadoop.repackaged.gcs.com.google.cloud.hadoop.gcsio.GoogleCloudStorageImpl: Ignoring exception of type GoogleJsonResponseException: verified object already exists with desired state.
Root-mean-square error = 0.868566272031658
+-----+-----+
|userId|movieId| rating|
+-----+-----+-----+
```

Root-mean-square error = 0.8685666272031658

userId	movieId	rating
471	3379	4.822564
471	8477	4.6659493
471	33649	4.5504856
471	102217	4.5333
471	92494	4.5333
471	33779	4.5333
471	171495	4.527984
471	7096	4.4821672
471	84273	4.4345856
471	117531	4.4345856
31	33649	5.0889573
31	3379	4.9877176
31	6086	4.85124
31	3200	4.813297
31	171495	4.79994
31	93988	4.786241
31	184245	4.7817674
31	84273	4.7817674
31	26073	4.7817674
31	7071	4.7817674

only showing top 20 rows

movieId	userId	rating	title	genres
67618	100	5.1201425	Strictly Sexual (...)	Comedy Drama Romance
3379	100	5.064743	On the Beach (1959)	Drama
42730	100	5.042285	Glory Road (2006)	Drama
33649	100	5.021657	Saving Face (2004)	Comedy Drama Romance
117531	100	4.9267745	Watermark (2014)	Documentary
7071	100	4.9267745	Woman Under the I...	Drama
184245	100	4.9267745	De platte jungle ...	Documentary
26073	100	4.9267745	Human Condition I...	Drama War
179135	100	4.9267745	Blue Planet II (2...	Documentary
84273	100	4.9267745	Zeitgeist: Moving...	Documentary

movieId	userId	rating	title	genres
1101	100	5.0	Top Gun (1986)	Action Romance
1958	100	5.0	Terms of Endearme...	Comedy Drama
2423	100	5.0	Christmas Vacatio...	Comedy
4041	100	5.0	Officer and a Gen...	Drama Romance
5620	100	5.0	Sweet Home Alabam...	Comedy Romance
368	100	4.5	Maverick (1994)	Adventure Comedy ...
934	100	4.5	Father of the Bri...	Comedy
539	100	4.5	Sleepless in Seat...	Comedy Drama Romance
16	100	4.5	Casino (1995)	Crime Drama
553	100	4.5	Tombstone (1993)	Action Drama Western

```
d24303a13656f17a345a38/driveroutput
jobUuid: 8494b30d-d291-35e9-a8b3-547accce96de
placement:
  clusterName: spark-cluster
  clusterUuid: a61a1a11-d3e2-46ad-8b23-ecfcaf287da7
pysparkJob:
  args:
    - --input_path_movies=gs://movie_recommendation_with_mllib_collaborative_filte/movies.csv
    - --input_path_ratings=gs://movie_recommendation_with_mllib_collaborative_filte/ratings.csv
  mainPythonFileUri: gs://movie_recommendation_with_mllib_collaborative_filte/Recommendation_Engine_MovieLens.py
reference:
  jobId: e4151b89e5d24303a13656f17a345a38
  projectId: cs570-big-data-424809
status:
  state: DONE
  stateStartTime: '2024-07-17T09:27:50.761112Z'
statusHistory:
- state: PENDING
  stateStartTime: '2024-07-17T09:00:46.828315Z'
- state: SETUP_DONE
  stateStartTime: '2024-07-17T09:00:46.862237Z'
- details: Agent reported job success
  state: RUNNING
  stateStartTime: '2024-07-17T09:00:47.054125Z'
yarnApplications:
- name: Recommendations
  progress: 1.0
  state: FINISHED
  trackingUrl: http://spark-cluster-m:8088/proxy/application_1721206458042_0002/
```